

MECH 6326 - Optimal Control and Dynamic Programming
Final Project Working Doc

Alyssa Vellucci and Jonas Wagner

May 4, 2023

Chapter 1

Simple System Model

1.1 System Definition

1.1.1 Assumptions:

- Movement: Single movement per turn
 - Deterministic
 - 1 square movement
 - move then action
- Actions: Single action per time step
 - Melee (hit check)(d6) - Short range
 - Ranged (hit check)(d8) - Longer range
 - Health Potion (d4 + 1)
 - Nothing
- Characters
 - 1 PC and 1 Monster
 - Identical Specs/modifiers
 - Same weapon (+2)
- Monster
 - Monster move in standard pattern
 - Monster cannot heal
- Infinite Time Horizon
- Infinite Battlefield and no Obstacles

1.1.2 Environment Definition

States

Let each character be associated with position and HP states. For position, let

$$x_{pc,p}, x_{mn,p} \in \mathcal{X}_p \subseteq \mathbb{Z}^2$$

describe the position on an infinite 2-d grid. For HP, let

$$x_{pc,hp}, x_{mn,hp} \in \mathcal{X}_{hp} \subseteq \mathbb{Z}_+ = \{0, 1, 2, \dots\}$$

describe the HP for each character.

Inputs

The inputs to the system consist of movement and actions impacting the position and hp states respectively. For movement, a deterministic input of

$$u_{pc,m}, u_{mn,m} \in \mathcal{U}_m = \{\text{Stop, N, E, S, W, NE, NW, SE, SW}\} \\ = \{(0,0), (-1,0), (+1,0), (0,-1), (0,+1), (-1,-1), (-1,+1), (+1,-1), (+1,+1)\} \quad (1.1)$$

For actions, the puts for the PC and monster are respectively

$$u_{pc,a}, u_{mn,a} \in \mathcal{U}_a = \{\text{Melee, Ranged, Heal, Nothing}\}$$

where each of the actions (except nothing) are stochastic and can either be represented as a function of input and noise terms or as markov chains.

For Melee and Ranged attacks, the character acts upon another character's HP where the impact on HP is as follows:

1. Ensure in range for either melee or ranged attack - otherwise can't attack.
2. "Roll" for success/fail - if fail then self-loop on opponent HP
3. "Roll" for effectiveness - opponent HP decreased by Weapon/self Modifiers (2) + d6/d8

The PC is also allowed to use a health potion which has a stochastic effect upon the player's health:

1. Ensure potion is available - otherwise can't heal
2. "Roll" for effectiveness - player's HP increased by health modifier (1) + d4

Stochastic Elements

When written in some forms the stochastic aspects of the system can be described as a noise signal consisting of the PC and monster dice rolls. The actions themselves can be modeled as either a function of is noise signal or as a derived markov chain acting directly upon the PC or Monster HP states.

The noise introduced by dice rolls are defined as for {d2, d4, d6, d8, d10, d20, d100} as

$$w_{pc,dn}, w_{mn,dn} \in \mathcal{W}_{dn} = \{1, 2, \dots, n\}$$

with $n = 2, 4, 6, 8, 10, 20, 100$ respectively where each outcome is equally likely. The PDF can be seen in shorthand with $P(w_{i,dn} = [1 \dots n]^T) = \frac{1}{n} [1 \dots 1]^T$.

Additionally, for the simplistic case, the success/failure and damage/heal amount can be determined independently. For the attack actions, success/failure is therefore dependent on the PC and monster's stats and a d20 roll, thus

$$w_{pc,sf}, w_{mn,sf} \in \mathcal{W}_{sf} \subseteq \{\text{"failure"} = 0, \text{"success"} = 1\}$$

and where $w_{pc,sf}$ and $w_{mn,sf}$ are directly calculated from PC and monster modifiers and their respective d20 rolls.

1.1.3 Problem Statement

For the simplistic case, let states at time-step k , be

$$x_k = \begin{bmatrix} x_{pc,p} \\ x_{mn,p} \\ x_{pc,hp} \\ x_{mn,hp} \\ x_{pc,potion} \end{bmatrix} \in \mathcal{X} = \mathcal{X}_p^2 \times \mathcal{X}_{hp}^2 \times \mathcal{X}_{potion} \subseteq \mathbb{Z}^4 \times \mathbb{Z}_+^2 \times \mathbb{Z}_+$$

where states and sets are defined as before and $x_{pc,potion} \in \mathcal{X}_{potion}$ is the number of potions available to the PC.

Let the inputs to the system be only the players inputs

$$u_k = \begin{bmatrix} u_{pc,m} \\ u_{pc,a} \end{bmatrix} \in \mathcal{U} = \mathcal{U}_p \times \mathcal{U}_a$$

The monster's inputs to the system will be incorporated as a deterministic input that and stochastic input that are closed-loop within the system and treated as part of the nonlinear aspects of the update function.

Let the noise signal w_k for each time-step be defined as a collection of the PC and Monster dice rolls,

$$w_k = \begin{bmatrix} w_{pc} \\ w_{mn} \end{bmatrix}, \quad w_i = \begin{bmatrix} w_{i,4} \\ w_{i,6} \\ w_{i,8} \\ w_{i,20} \end{bmatrix} \forall_{i=pc,mn}$$

, and the associated success/fail noise $\forall_{i=pc,mn} \forall_{j=mn,pc} \forall_{l=dex,str}$ are derived as

$$w_{i,sf} = \begin{cases} 0 & (a_{i,l} + w_{i,d20} \leq a_{j,ac} \text{ OR } w_{i,d20} = 1) \text{ AND } w_{i,d20} \neq 1 \\ 1 & (a_{i,l} + w_{i,d20} > a_{j,ac} \text{ OR } w_{i,d20} = 20) \text{ AND } w_{i,d20} \neq 1 \end{cases}$$

where the modifiers are held constant as $a_{i,l} = a_{i,dex} = a_{i,str} = 5$ and $a_{j,ac} = 15$.

This results in probabilities of $P(w_{i,sf} = 0) = (P(w_{i,d20} = 1) + P(5 + w_{i,d20} < 15)) - (P(w_{i,d20} = 20)) = 1/20 + 9/20 - 1/20 = 9/20$ and $P(w_{i,sf} = 1) = (P(w_{i,d20} = 20) + P(5 + w_{i,d20} \geq 15)) - (P(w_{i,d20} = 1)) = 1/20 + 11/20 - 1/20 = 11/20$. In shorthand, $P(w_{i,sf} = \begin{bmatrix} 0 \\ 1 \end{bmatrix}) = \begin{bmatrix} 0.45 \\ 0.55 \end{bmatrix}$.

The evolution of the very simple system can be described as Markov chains or by a nonlinear update function:

$$x_{k+1} = f(x_k, u_k, w_k) = \begin{bmatrix} x_{pc,p} + f_{pc,m}(u_k) \\ x_{mn,p} + f_{mn,m}(x_k) \\ \begin{bmatrix} x_{pc,hp} \\ x_{mn,hp} \\ x_{pc,potion} \end{bmatrix} + f_{pc,a}(x_k, u_k, w_k) + f_{mn,a}(x_k, w_k) \end{bmatrix}$$

where the associated functions update states as follows:

- The players deterministic movement input: $f_{pc,m}(u_k) = u_{pc,m}$
- The monsters state-dependent movement:

$$f_{mn,m}(x_k) = \text{direction}(x_{pc,p} - x_{mn,p})$$

where $\text{direction}()$ is calculated as the closest cardinal direction that heads towards the player. (In matlab: `round(k_mn_speed * normalize(x_pc_p - x_mn_p))`)

- The players action:

$$f_{pc,a} = \begin{cases} \begin{bmatrix} 0 & -w_{pc,sf}(a_{pc,wpn,m} + w_{pc,d6}) & 0 \end{bmatrix}^T & u_{pc,a} = \text{Melee AND } \|x_{pc,p} - x_{mn,p}\|_2 \leq a_{pc, rng, m} \\ \begin{bmatrix} 0 & -w_{pc,sf}(a_{pc,wpn,r} + w_{pc,d8}) & 0 \end{bmatrix}^T & u_{pc,a} = \text{Ranged AND } a_{pc, rng, m} < \|x_{pc,p} - x_{mn,p}\|_2 \leq a_{pc, rng, r} \\ \begin{bmatrix} a_{pc,potion} + w_{pc,d4} & 0 & -1 \end{bmatrix}^T & u_{pc,a} = \text{Heal AND } x_{pc,potion} \geq 1 \\ \begin{bmatrix} 0 & 0 & 0 \end{bmatrix}^T & u_{pc,a} = \text{Nothing OR Otherwise} \end{cases}$$

- The monsters state dependent action:

$$f_{mn,a} = \begin{cases} \begin{bmatrix} -w_{mn,sf}(a_{mn,wpn,m} + w_{mn,d6}) & 0 & 0 \end{bmatrix}^T & \|x_{pc,p} - x_{mn,p}\|_2 \leq a_{mn, rng, m} \\ \begin{bmatrix} -w_{mn,sf}(a_{mn,wpn,r} + w_{mn,d8}) & 0 & 0 \end{bmatrix}^T & a_{mn, rng, m} < \|x_{pc,p} - x_{mn,p}\|_2 \leq a_{mn, rng, r} \\ \begin{bmatrix} 0 & 0 & 0 \end{bmatrix}^T & a_{mn, rng, r} < \|x_{pc,p} - x_{mn,p}\|_1 \end{cases}$$

where weapons/potion modifiers are taken as $a_{pc,wpn,m} = a_{mn,wpn,m} = 0$, $a_{pc,wpn,r} = a_{mn,wpn,r} = 0$, and $a_{pc,potion} = 1$ and ranges are set as $a_{pc,rng,m} = a_{mn,rng,m} = 2$ and $a_{pc,rng,r} = a_{mn,rng,r} = 5$

The objective function is defined to minimize the monsters HP and maximize the PC HP while incentivising monster death and strongly disincentivising PC death.

1.1.4 Markov Implementation

TODO: Simplify the Markov implementation section. It may be easier to incorporate the relative P definitions alongside the stochastic function equivalent and then have only the absolute markov definitions in the MATLAB implementation.

For each of the stochastic functions defined in the update functions from the problem definition a markov chain can be described to better represent the changes between different states. Considering that only the actions are stochastic, the chains can be restricted to only the PC and Monster hp where the action input decides which of the chains are selected.

A few additional assumptions must be made to restrict the problem to a markov chain implementation:

- The monster and pc hp will be restricted to a finite range, $x_{pc,hp}, x_{mn,hp} \in \mathcal{X}_{hp} = \{0, 1, \dots, n_{hp,max}\}$
- The ensuring of a potion will be a restriction upon the pc input/result in a nothing action if no potion is available.
- The deterministic movement will allow for monster action selection within closed-loop, however this splits the actual markov chain in two.
- The actual implimentation will keep the pc and monster hp independent of one another with actions being applied to the appropriate hp according to the action; however a proper markov chain would incorporate both into the states.

The states of the markov chain will be artificially independent $x_{pc,k}, x_{mn,k} = P(x_{pc,hp}, x_{mn,hp}) \in [0, 1]^{n_{hp,max}+1}$. The transition matrices will be defined as $[p_{i,j}] = P(x_{k+1} = i, x_k = j)$. This can be thought of as each column being the probabilistic function results for $f(x_k, w_k), x_k = j, w_k = P[w_k]$.

PC Actions

First for the pc actions, which will be done independently from the monsters action which is then applied after according to the position.

When $u_k = \text{Melee}$, the $f(x_k, w_k)$ will act upon $x_{mn,hp}$ with uncertainty $w_k \in \mathcal{W}_{d6}$, resulting in $f(x_{k+1} = i : x_k = j) = j - w_{pc,melee,sf} * (2 + w_{pc,d6}) = j - P(w_{pc,melee,sf} = [0]) \times (2 + w_{pc,d6} P(w_{pc,d6} = [1 \dots 6])) = j - [0.45] \times (2 + w_{pc,d6} [1/6 \dots 1/6]^T)$ resulting in $P(x_{k+1} = j) = 0.45$, $P(x_{k+1} = j - 2 - 1) = \dots = P(x_{k+1} = j - 2 - 6) = \frac{0.55}{6} = \frac{11}{120} \approx 0.01967$. This is incorporated into the markov chain as the column j has

$$\begin{array}{c} i = 1 \\ \vdots \\ i = j - 2 - 7 \\ i = j - 2 - 6 \\ \vdots \\ i = j - 2 - 1 \\ i = j - 2 \\ i = j - 1 \\ i = j \\ i = j + 1 \\ \vdots \end{array} \begin{bmatrix} 0 \\ \vdots \\ 0 \\ 0.019677 \\ \vdots \\ 0.01967 \\ 0 \\ 0 \\ 0 \\ 0.45 \\ 0 \\ \vdots \end{bmatrix}$$

and the first few columns will have all the probabilities of negative values summed into the zero row. The markov chain transition matrix will be denoted as $P_{pc,melee}$.

Improve notation of this terrible thing...

Notation terrible

confirm this direction... is this P or P^T?

Essentially the same definition is true with $u_k = \text{ranged}$ but with $w_{i,d8}$ and resulting in $P(x_{k+1} = j) = 0.45$, $P(x_{k+1} = j - 2 - 1) = \dots = P(x_{k+1} = j - 2 - 8) = \frac{0.55}{8} = \frac{11}{160} = 0.06875$ and denoted with markov chain transition matrix $P_{pc,ranged}$.

Similarly for $u_k = \text{heal}$ but with $w_{i,d4}$ and not having the w_{sf} , meaning $P(x_{k+1} = j + 1 + 1) = \dots = P(x_{k+1} = j + 1 + 4) = 1/4 = 0.25$ and denoted with markov chain transition matrix $P_{pc,heal}$.

Finally, for $u_k = \text{nothing}$ the transition matrix is defined as the identity matrix: $P_{pc,nothing} = \mathbf{I}_{n_{hp,max}}$

Monster Movement and Actions

For the simplistic case the movement is done first deterministically and then the monster actions are performed using the markov chains. Selection between the actions would be done deterministically based on position after movement and then the action is the same transition matrix for melee $P_{mn,melee}$, ranged $P_{mn,ranged}$, and nothing $P_{mn,nothing}$ are equivalent to the PC's transition matrices as $P_{pc,melee}$, $P_{pc,ranged}$, and $P_{pc,nothing}$ respectively.

The movement is done as described in the problem definition: $x_{k+1,mn,p} = x_{k,mn,p} + \text{direction}(x_{k+1,pc,p} - x_{k,mn,p})$.

The deterministic action input results in the monsters action being applied

$$P_{mn} = \begin{cases} P_{mn,melee} & \|x_{pc,p} - x_{mn,p}\|_1 \leq \text{Melee Range} \\ P_{mn,ranged} & \text{Melee Range} < \|x_{pc,p} - x_{mn,p}\|_1 \leq \text{Ranged Range} \\ P_{mn,nothing} & \text{Ranged Range} < \|x_{pc,p} - x_{mn,p}\|_1 \end{cases}$$

where each will act upon the PC hp.

Closed-loop Markov Chain

The actual markov chains associated with closed-loop will be selected directly based on state and player input by selecting the appropriate markov chains.

Since the assumption that u_k will be limited to those feasible by the PC, the following described as follows:

First, PC movement: $x_{k+1,pc,p} = x_{k,pc,p} + u_{pc,m}$.

Second, PC action:

$$\begin{bmatrix} x_{k+1,pc,hp} \\ x_{k+1,mn,hp} \end{bmatrix} = \begin{cases} \begin{bmatrix} x_{k,pc,hp} \\ x_{k,mn,hp}^T P_{pc,melee} \end{bmatrix} & u_k = \text{melee} \\ \begin{bmatrix} x_{k,pc,hp} \\ x_{k,mn,hp}^T P_{pc,ranged} \end{bmatrix} & u_k = \text{ranged} \\ \begin{bmatrix} x_{k,pc,hp}^T P_{pc,heal} \\ x_{k,mn,hp} \end{bmatrix} & u_k = \text{heal} \\ \begin{bmatrix} x_{k,pc,hp} \\ x_{k,mn,hp} \end{bmatrix} & u_k = \text{nothing} \end{cases}$$

Third, monster movement: $x_{k+1,mn,p} = x_{k,mn,p} + \text{direction}(x_{k+1,pc,p} - x_{k,mn,p})$.

Finally, monster action:

$$x_{k+1,pc,hp} = \begin{cases} x_{k,pc,hp}^T P_{mn,melee} & \|x_{k+1,pc,p} - x_{k+1,mn,p}\|_1 \leq \text{Melee Range} \\ x_{k,pc,hp}^T P_{mn,ranged} & \text{Melee Range} < \|x_{k+1,pc,p} - x_{k+1,mn,p}\|_1 \leq \text{Ranged Range} \\ x_{k,pc,hp} & \text{Ranged Range} < \|x_{k+1,pc,p} - x_{k+1,mn,p}\|_1 \end{cases}$$

It is possible to incorporate this into a single set of markov chains in higher dimensions; however, this sequential definition seems to make the most sense instead of defining a finite number of position states and closing the loop entirely.

being able to move after certainly would add a lot more capability to the player's actions... however this would absolutely complicate it more...

1.2 MATLAB implimentation

1.2.1 Additional Simplifications

Due to the nature of the problem being of such high dimension, the following additional simplifications will be done:

- The policy will be only based on relative distances between the PC and Monster. (i.e.)

$$x_{k,pos} = x_{pc,p} - x_{mn,p} = \begin{bmatrix} x_{pc,p,x} - x_{mn,p,x} \\ x_{pc,p,y} - x_{mn,p,y} \end{bmatrix}$$

This is mainly to reduce the number of dimensions from 4 to 2 and ensure the finite position is useful.

- The objective function will be only based on the relative *HP* between PC and monster, $x_{pc,hp} - x_{mn,hp}$. Specifically, the stage-cost function will be to maximize this relative *HP*:

$$g_k(x_{pc}, x_{mn,hp}) = -(x_{pc,hp} - x_{mn,hp})$$

- A value iteration method will be used (first attempted in Infinite-Horizon... but might be a finite horizon instead).
- When determining the policy, the cost-function will be augmented to discourage moving too far away and a strong discouragement of losing (or encouragement of winning) will be included.

1.2.2 Optimal Policy

Update Probabilities

The probabilities fore each state update (essentially a large-multi dimensional markov-chain) will be computed ahead of time to make the policy optimization more computationally efficient. Specifically the closed-loop markov chains will be used to determine the probabilities of future states $\mathbf{P}[x_{k+1}]$ that can then be used for the expected next time-step cost computation:

$$\mathbf{E}[J(f(x_k, u_k, w_k))] = \mathbf{P}[x_{k+1}]^T J_{k+1}(x_{k+1})$$

Value Iteration

The optimal policy will be determined for every relative position and the current HPs. These states will be treated as within a 4-D grid which are looped through (potentially in parallel for computation purposes... or maybe a vector version can be determined...) and the optimal policy will be updated each iteration.

This could probably instead be done with a much better set of matrix/vector operations but that isn't yet figured out...

this could be explained much better... but essentially use the code from previous examples...